

25 YEARS ANNIVERSARY
SICT

HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY



HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Feature engineering

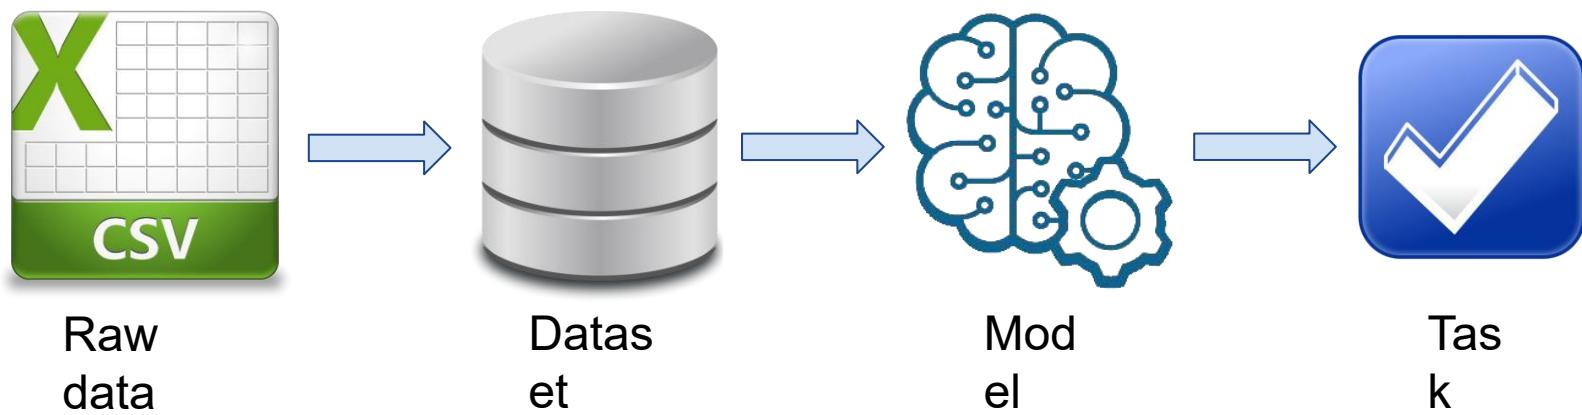
Feature engineering

- "Feature engineering is the process of transforming raw data into features that better represent the underlying problem to the predictive models, resulting in improved model accuracy on unseen data." – Jason Brownlee

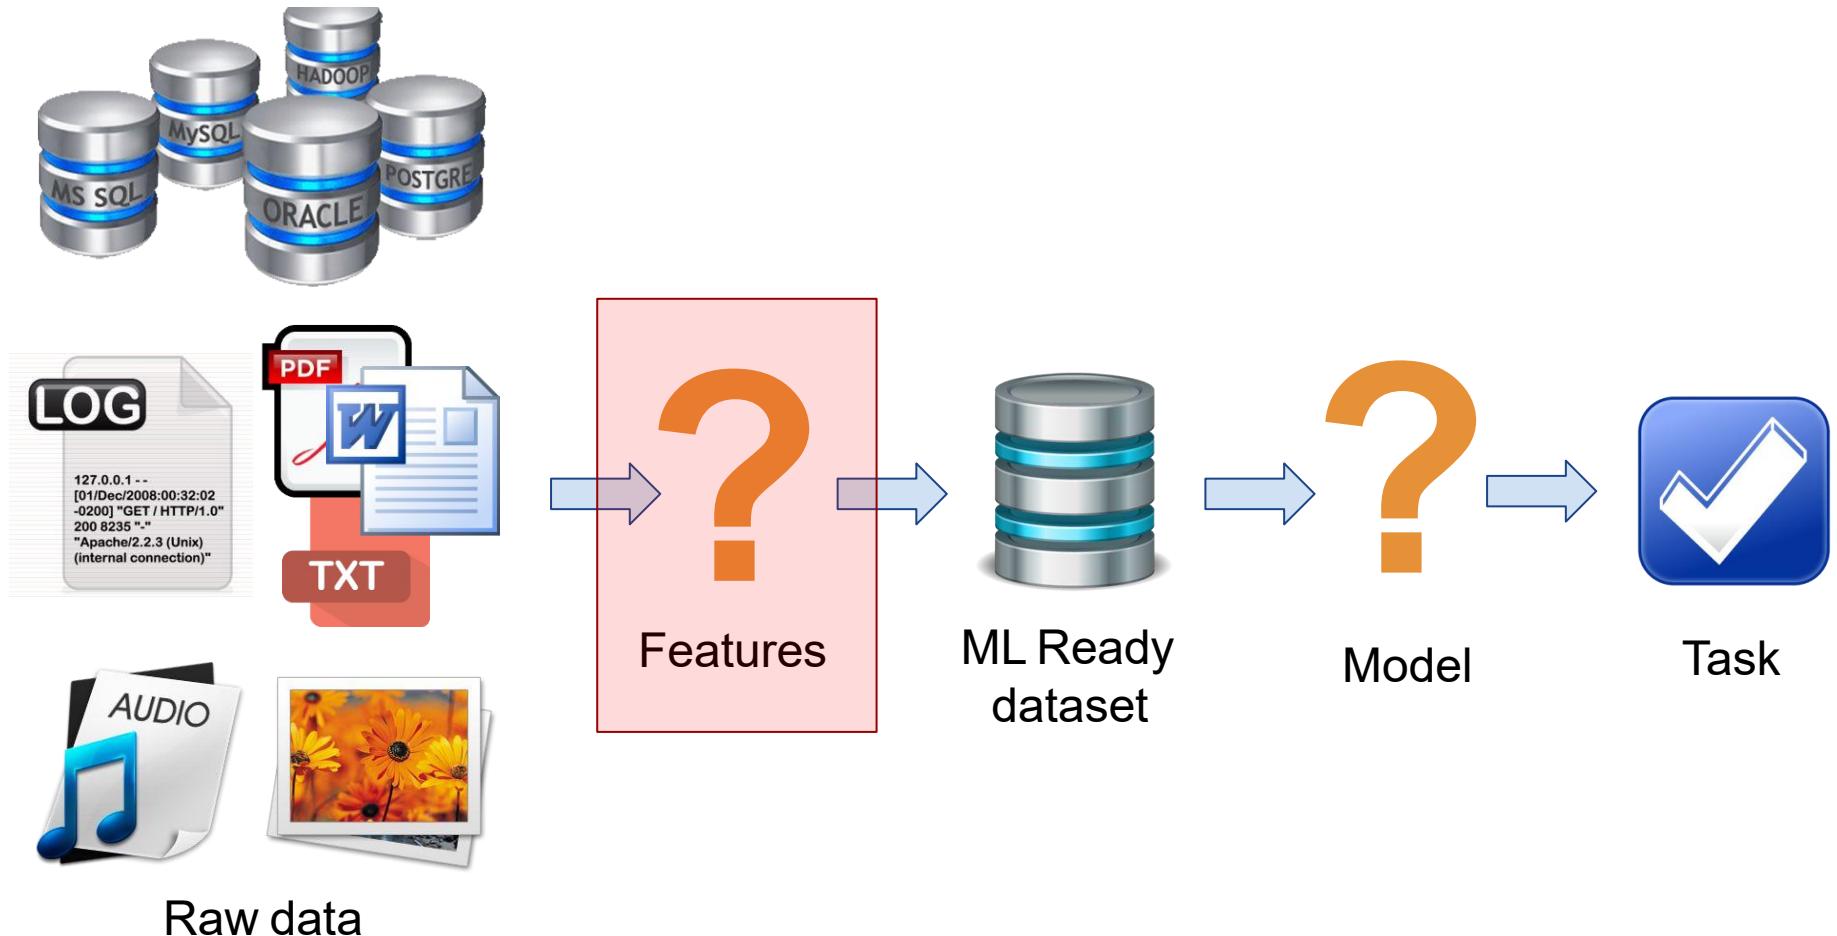
Feature engineering

- “Coming up with features is difficult, time-consuming, requires expert knowledge. 'Applied machine learning' is basically feature engineering.” – Andrew Ng

The dream ...



... The Reality

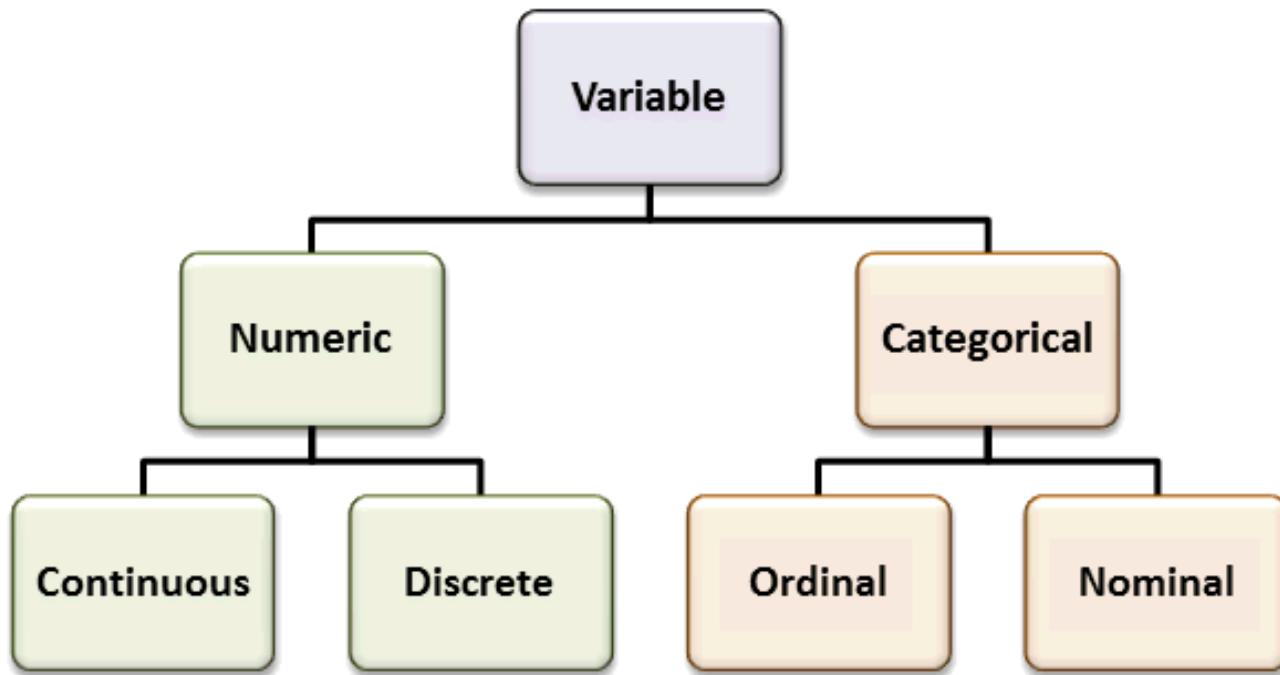


Feature engineering toolbox

- Just kidding :)



Variable data types



Number variables

Binarization

- Counts can quickly accumulate without bound
- convert them into binary values (0, 1) to indicate presence

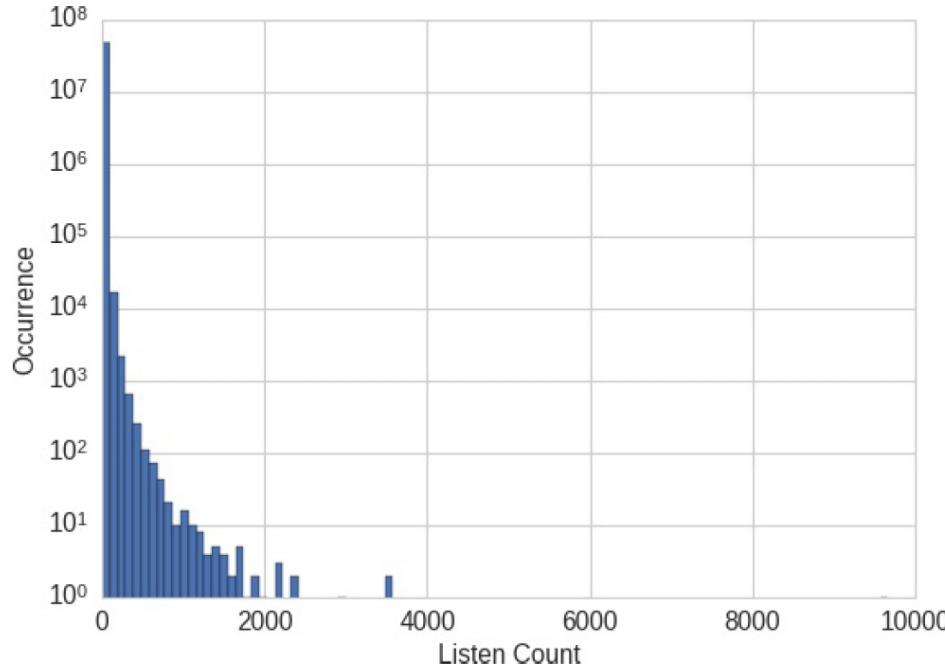
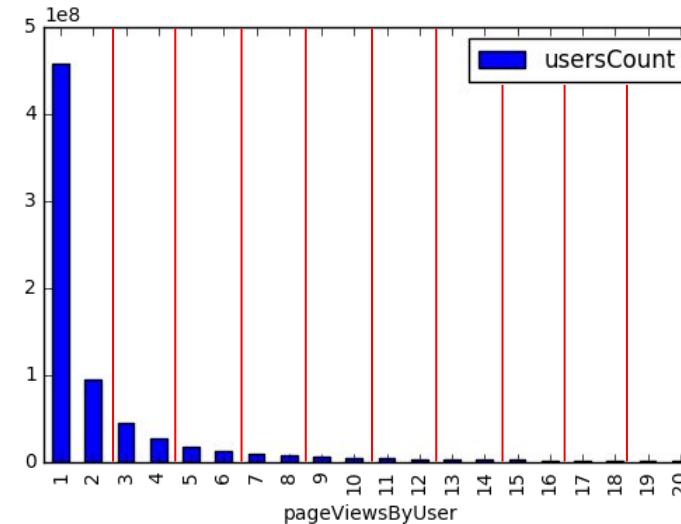


Figure 2-3. Histogram of listen counts in the Taste Profile subset of the *Million Song Dataset*—
note that the y-axis is on a log scale

Quantization or Binning

- Group the counts into bins
- Maps a continuous number to a discrete one
- Bin size
 - Fixed-width binning
 - Eg.
 - 0–12 years old
 - 12–17 years old
 - 18–24 years old
 - 25–34 years old
 - Adaptive-width binning



Equal Width Binning

- divides the continuous variable into several categories having bins or range of the **same width**

Categories : $[min, min + w - 1], [min + w, min + 2 * w - 1], [min + 2 * w, min + 3 * w - 1] \dots [min + (x - 1) * w, max]$

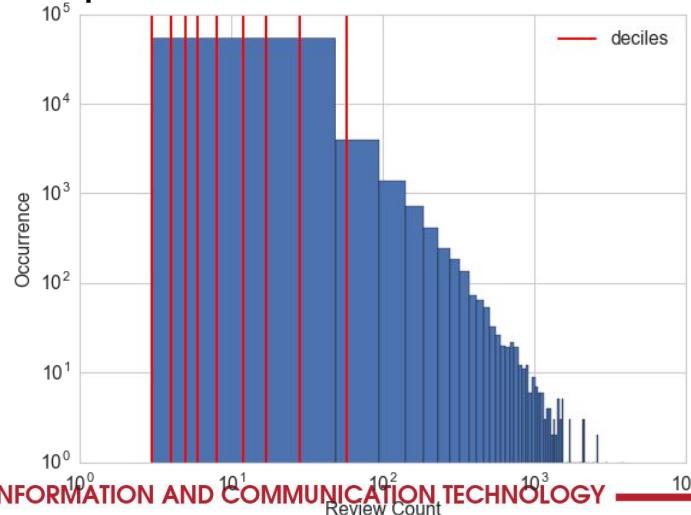
$$w = \left\lfloor \frac{max - min}{x} \right\rfloor$$

- Pros
 - easy to compute
- Cons
 - large gaps in the counts
 - many empty bins with no data

AGE	AGE_bins
10	[10, 21]
15	[10, 21]
16	[10, 21]
18	[10, 21]
20	[10, 21]
30	[22, 33]
35	[34, 45]
42	[34, 45]
48	[46, 55]
50	[46, 55]
52	[46, 55]
55	[46, 55]

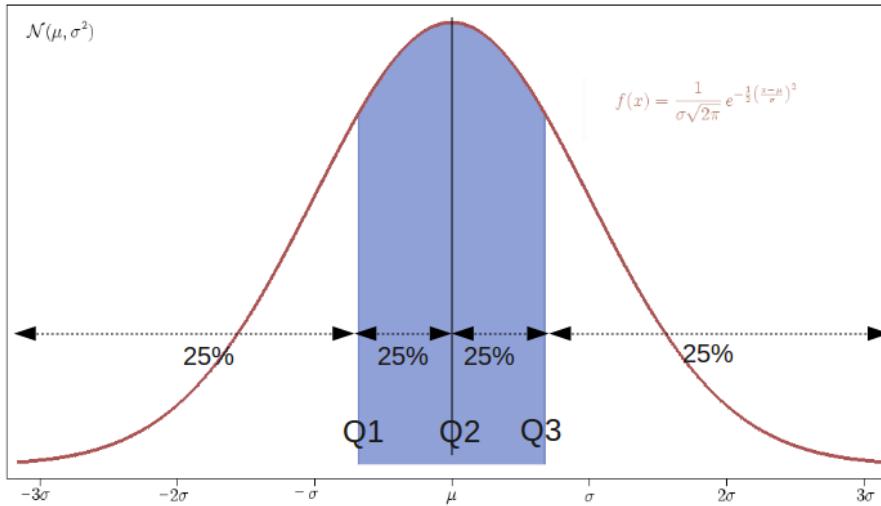
Adaptive-width binning

- Equal frequency binning
 - Quantiles: values that divide the data into equal portions (continuous intervals with equal probabilities)
 - Some q-quantiles have special names
 - The only 2-quantile is called the median
 - The 4-quantiles are called quartiles → Q
 - The 6-quantiles are called sextiles → S
 - The 8-quantiles are called octiles
 - The 10-quantiles are called deciles → D



AGE	AGE_bins
10	[10, 16]
15	[10, 16]
16	[10, 16]
18	[17, 30]
20	[17, 30]
30	[17, 30]
35	[31, 48]
42	[31, 48]
48	[31, 48]
50	[49, 55]
52	[49, 55]
55	[49, 55]

Example: quartiles



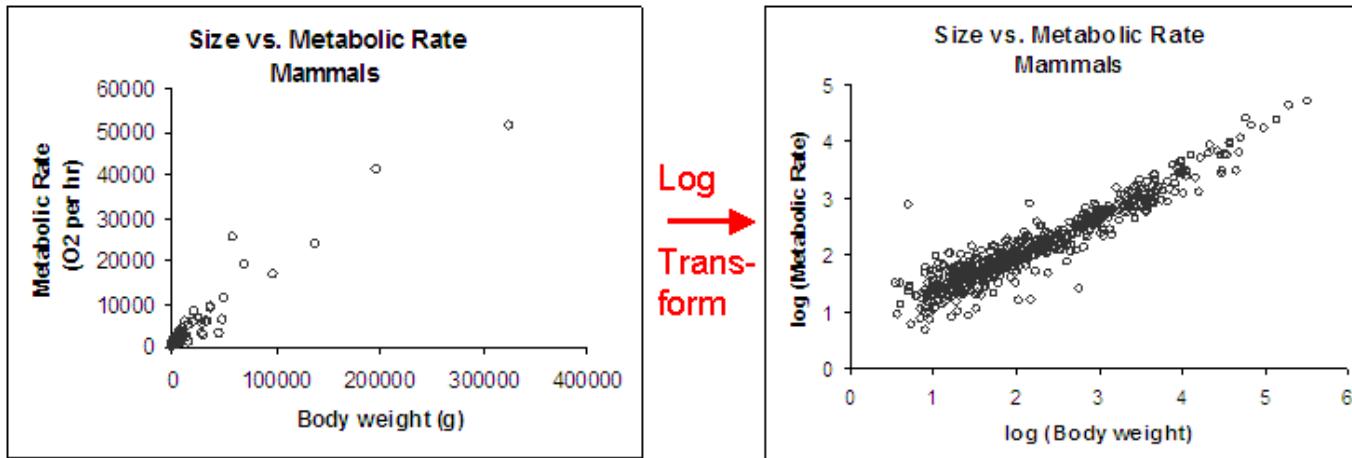
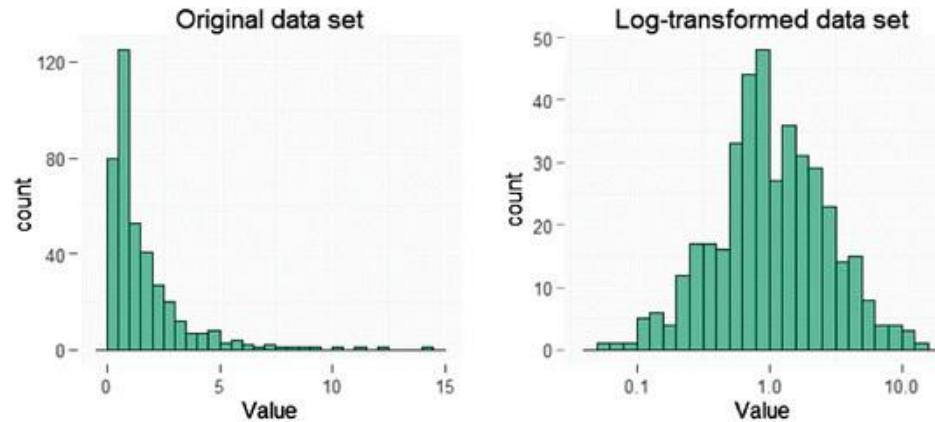
```
>>> import pandas as pd

# Map the counts to quartiles
>>> pd.qcut(large_counts, 4, labels=False)
array([1, 2, 3, 0, 0, 1, 1, 2, 2, 3, 3, 0, 0, 0, 2, 1, 0, 3], dtype=int64)

# Compute the quantiles themselves
>>> large_counts_series = pd.Series(large_counts)
>>> large_counts_series.quantile([0.25, 0.5, 0.75])
0.25    122.0
0.50    926.0
0.75   8286.0
dtype: float64
```

Log Transformation

- Original number = x
- Transformed number
 $x'=\log_{10}(x)$
- Backtransformed number =
 $10^{x'}$



Box-Cox transformation

$$\tilde{x} = \begin{cases} \frac{x^\lambda - 1}{\lambda} & \text{if } \lambda \neq 0, \\ \ln(x) & \text{if } \lambda = 0. \end{cases}$$

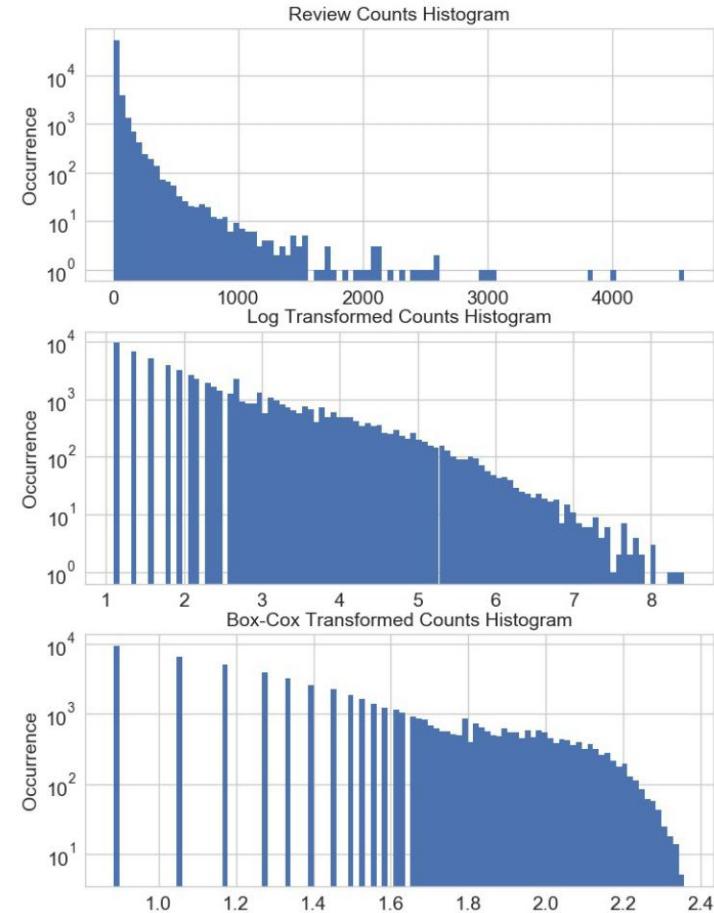


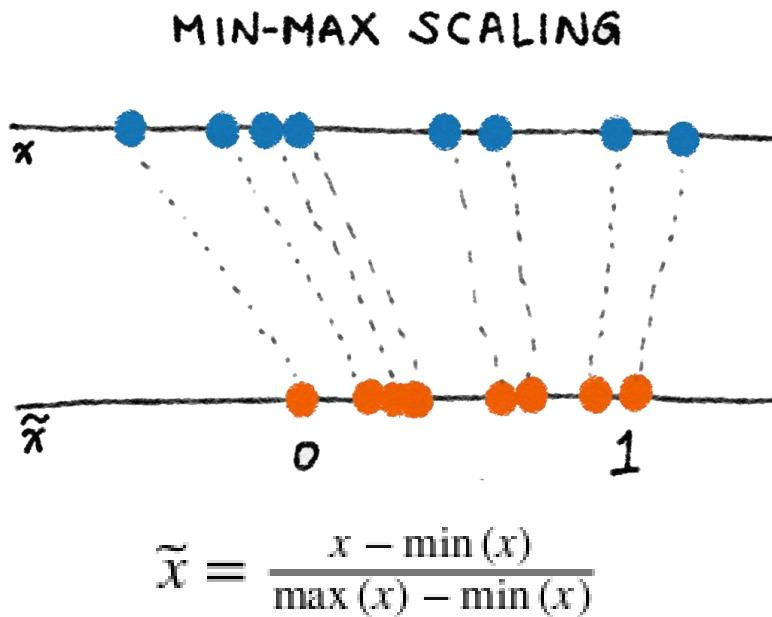
Figure 2-13. Box-Cox transformation of Yelp business review counts (bottom), compared to original (top) and log transformed (middle) histograms

Feature Scaling (Normalization)

- Models that are smooth functions of the input, such as linear regression, logistic regression are affected by the scale of the input
- Feature scaling or normalization changes the scale of the features

Min-max scaling

- Squeezes (or stretches) all values within the range of [0, 1] to add robustness to very small standard deviations and preserving zeros for sparse data.

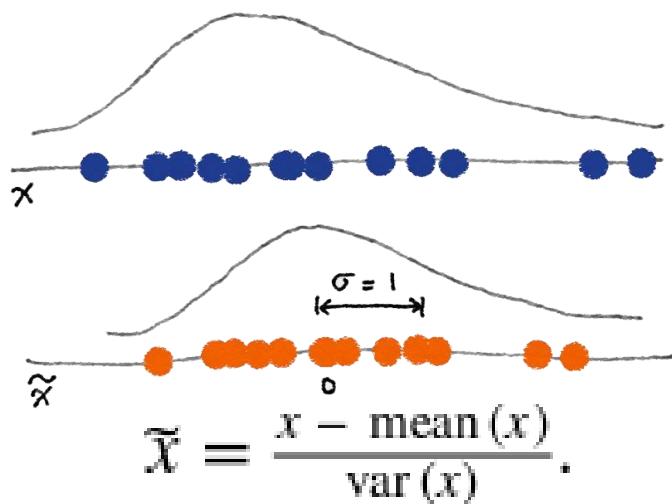


```
• >>> from sklearn import preprocessing  
• >>> X_train = np.array([[ 1., -1.,  2.,  
...                   [ 2.,  0.,  0.],  
...                   [ 0., 1., -1.]])  
• ...  
• >>> min_max_scaler =  
preprocessing.MinMaxScaler()  
• >>> X_train_minmax =  
min_max_scaler.fit_transform(X_train)  
array([[ 0.5      ,  0.       ,  1.       ],  
       [ 1.       ,  0.5     ,  0.33333333],  
       [ 0.       ,  1.       ,  0.       ]])
```

Standard (Z) Scaling

After Standardization, a feature has mean of 0 and variance of 1 (assumption of many learning algorithms)

STANDARDIZATION



```
>>> from sklearn import preprocessing
>>> import numpy as np
>>> X = np.array([[1., -1., 2.],
   ...             [2., 0., 0.],
   ...             [0., 1., -1.]])
>>> X_scaled = preprocessing.scale(X)
>>> X_scaled
array([[ 0. ..., -1.22...,  1.33...],
       [ 1.22...,  0. ..., -0.26...],
       [-1.22...,  1.22..., -1.06...]])
>> X_scaled.mean(axis=0)
array([ 0.,  0.,  0.])
>>> X_scaled.std(axis=0)
array([ 1.,  1.,  1.])
```

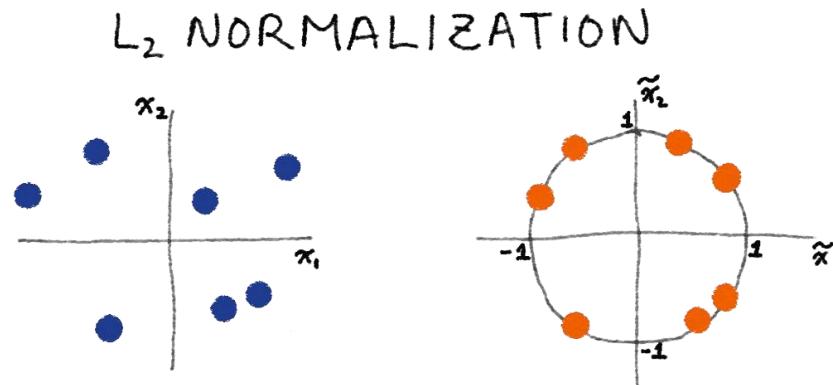
Standardization with scikit-learn

$\| \cdot \|_2$ Normalization

- also known as the Euclidean norm
- measures the length of the vector in coordinate space
- scale the values so that if they were all squared and summed, the value would be 1

$$\tilde{x} = \frac{x}{\|x\|_2} \quad \|x\|_2 = \sqrt{x_1^2 + x_2^2 + \dots + x_m^2}$$

```
from pandas import read_csv
from numpy import set_printoptions
from sklearn.preprocessing import Normalizer
path = r'./pima-indians-diabetes.csv'
names = ['preg', 'test', 'mass', 'pedi', 'age', 'class']
dataframe = read_csv(path, names=names)
array = dataframe.values
Data_normalizer = Normalizer(norm='l2').fit(array)
Data_normalized = Data_normalizer.transform(array)
```



Categorical Variables

Categorical Features

- Nearly always need some treatment to be suitable for models
- High cardinality can create very sparse data
- Difficult to impute missing
- Examples
 - Platform: [“desktop”, “tablet”, “mobile”]
 - Document_ID or User_ID: [121545, 64845, 121545]

Label Encoding

- transform categorical variables into numerical variables by assigning a numerical value to each of the categories

[male, female]	[0, 1]
[blue, green, red, black]	[0, 1, 2, 3]
[10, 21], [22, 33], [34, 45], [46, 55]	[0, 1, 2, 3]

LabelCount encoding

- Rank categorical variables by count in train set
- Useful for both linear and non-linear algorithms (eg: decision trees)
- Not sensitive to outliers
- Won't give same encoding to different variables

ad_id	clicks	ad_rank
54345	35387	1
423654	18339	2
98799	12352	3
68655	9430	4
123646	8232	5

Ordinal encoding

- transform an original categorical variable to a numerical variable by **ensuring the ordinal nature of the variables is sustained**

[male, female]	[0, 1]
[10, 21], [22, 33], [34, 45], [46, 55]	[0, 1, 2, 3]
[cold, warm, hot]	[0, 1, 2]
[poor, fair, good, very good, excellent]	[0, 1, 2, 3, 4]

Frequency encoding

- transform an original categorical variable to a numerical variable by considering the frequency distribution of the data

Column	Freq_Encoding
red	5
green	3
red	5
green	3
blue	4
red	5
red	5
blue	4
red	5
blue	4
blue	4
green	3

One hot encoding

- creates k different columns each for a category and replaces one column with 1 rest of the columns is 0

Column	red	green	blue
red	1	0	0
green	0	1	0
red	1	0	0
green	0	1	0
blue	0	0	1
red	1	0	0
red	1	0	0
blue	0	0	1
red	1	0	0
blue	0	0	1
blue	0	0	1
green	0	1	0

Target Mean encoding

- one of the best techniques
- replace the categorical variable with the mean of its corresponding target variable
- **Steps for mean encoding**
 - For each category
 - Calculate aggregated sum (= a)
 - Calculate aggregated total count (= b)
 - Numerical value for that category = a/b

Column	Target	Target Mean	Target Mean (numerical value)
red	1	3/5	0.6
green	1	2/3	0.67
red	0	3/5	0.6
green	0	2/3	0.67
blue	1	2/4	0.5
red	0	3/5	0.6
red	1	3/5	0.6
blue	0	2/4	0.5
red	1	3/5	0.6
blue	0	2/4	0.5
blue	1	2/4	0.5
green	1	2/3	0.67

Feature Hashing

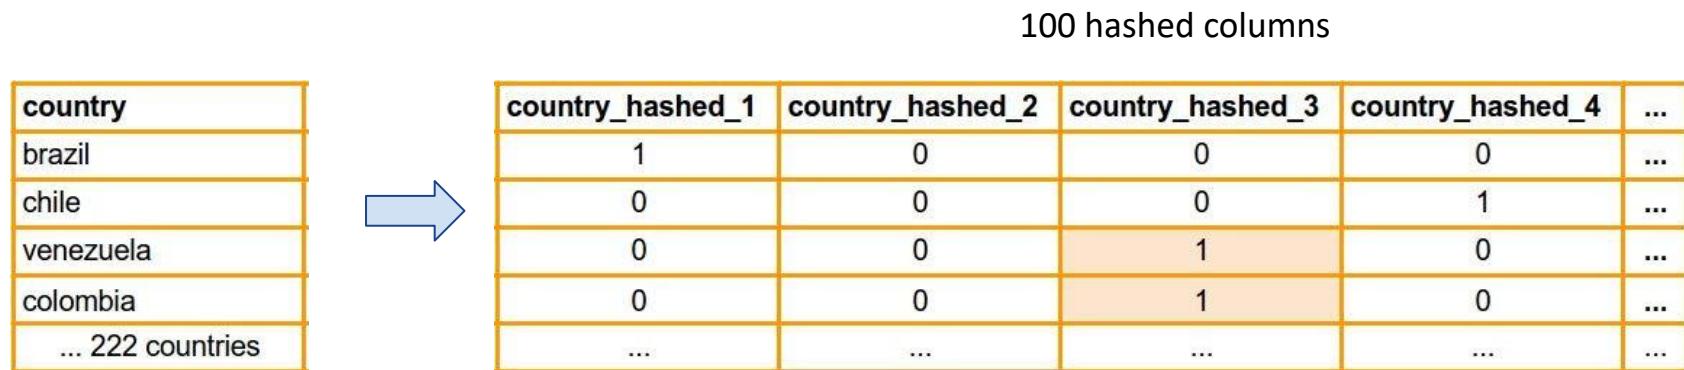
- Dealing with Large Categorical Variables

categorical_feature	unique_values
landing_page_document_id	636482
ad_id	418295
ad_document_id	143856
content_entities	52439
advertiser	2052
publisher	830
country_state	1892

Some large categorical features from Outbrain Click Prediction competition

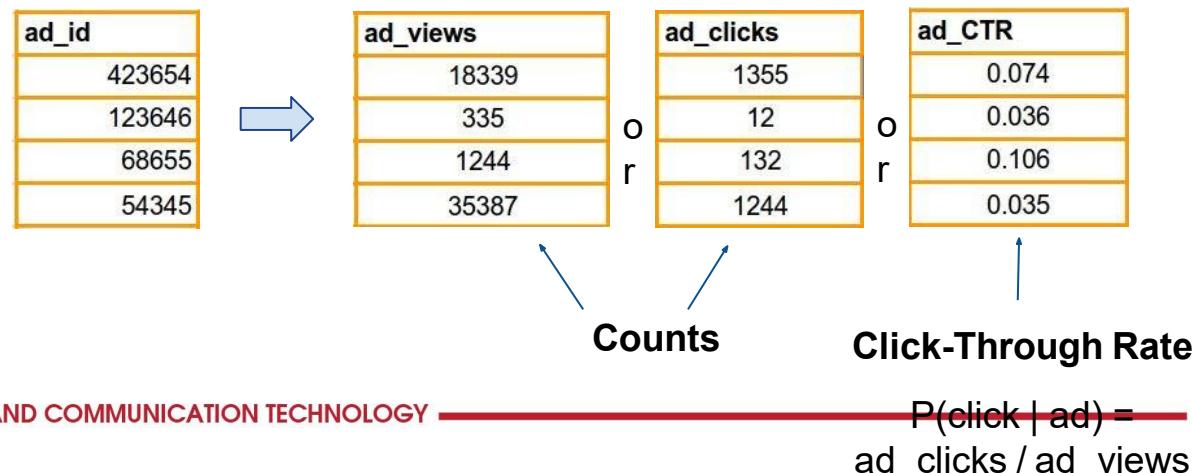
Feature hashing [2]

- Hashes categorical values into vectors with fixed-length.
- Lower sparsity and higher compression compared to one hot encoding
- Deals with new and rare categorical values (eg: new user-agents)
- May introduce collisions



Bin-counting

- Instead of using the value of the categorical variable as the feature, **we compute the association statistics between that value and the target that we wish to predict**
- Useful for both linear and non-linear algorithms
- May give collisions (same encoding for different categories)
- Be careful about leakage
- Strategies
 - Count
 - Average CTR



Text Data

Natural Language Processing

- **Cleaning**
- Lowercasing
- Convert accented characters
- Removing non-alphanumeric
- Repairing
- **Tokenizing**
- Encode punctuation marks
- Tokenize
- N-Grams
- Skip-grams
- Char-grams
- Affixes
- **Removing**
- Stopwords
- Rare words
- Common words
- **Roots**
- Spelling correction
- Chop
- Stem
- Lemmatize
- **Enrich**
- Entity Insertion / Extraction
- Parse Trees
- Reading Level

Text vectorization

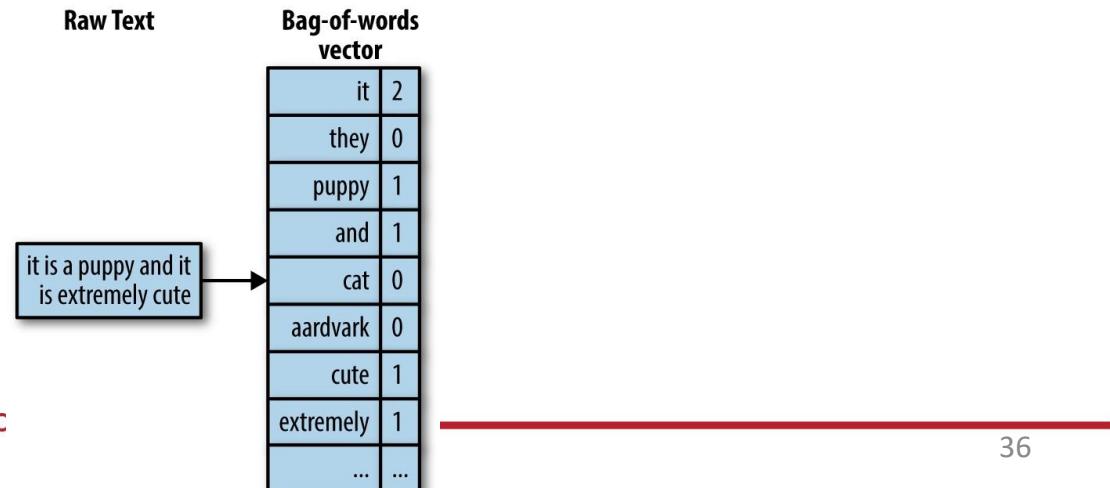
- Represent each document as a feature vector in the vector space, where each position represents a word (token) and the contained value is its relevance in the document.
 - BoW (Bag of words)
 - TF-IDF (Term Frequency - Inverse Document Frequency)
 - Embeddings (eg. Word2Vec, Glove)
 - Topic models (e.g LDA)

	linux	modern	the	system	steering	petrol
D1	3	4	3	0	2	0
D2	4	3	4	1	0	1
D3	1	0	4	1	0	1
D4	0	1	3	3	3	4

Document Term Matrix - Bag of Words

Bag-of-Words

- Input
 - “Customer reviews build something known as social proof, a phenomenon that states people are influenced by those around them. This might include friends and family, industry experts and influencers, or even internet strangers.”
- Output
 - a text document is converted into a “flat” vector of counts
 - doesn’t contain any of the original textual structures
 - *John is quicker than Mary and Mary is quicker than John have the same vectors*

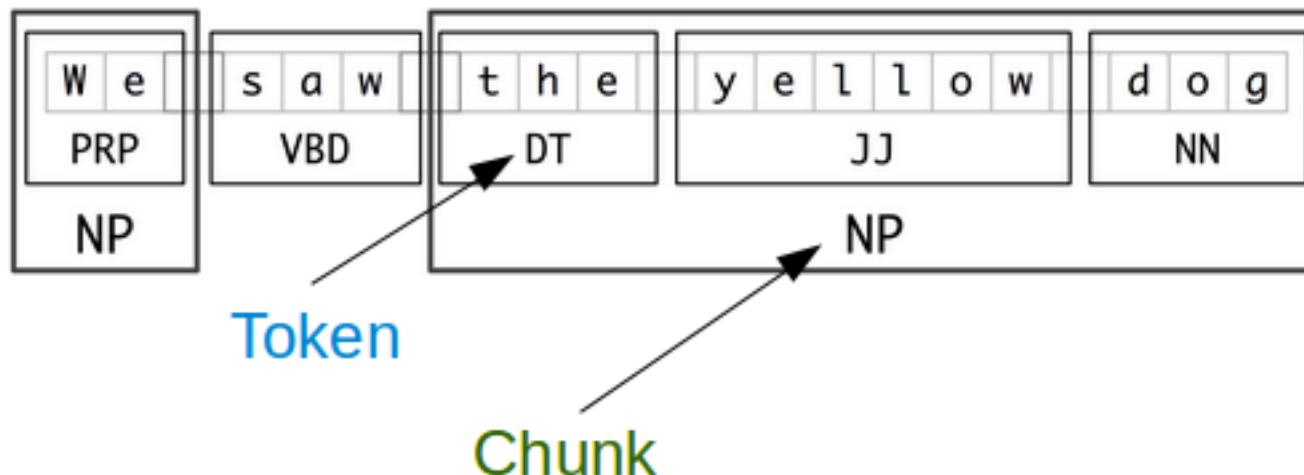


Bag-of-n-Grams

- a natural extension of bag-of-words (a word is essentially an unigram)
- bag-of-n-grams representation can be more informative
 - n-grams retain more of the original sequence structure
- Cons
 - bag-of-n-grams is a much bigger and sparser feature space

From Words to n-Grams to Phrases

- **Tokenization** is the process of **tokenizing** or splitting a string, text into a list of tokens.
- **Chunking** a **sentence** refers to breaking/dividing a **sentence** into parts of words such as word groups and verb groups.



Document frequency

- Rare terms are more informative than frequent terms
 - Recall stop words
- Consider a term in the query that is rare in the collection (e.g., *arachnocentric*)
- A document containing this term is very likely to be relevant to the query *arachnocentric*
 - → We want a high weight for rare terms like *arachnocentric*.

Tf-idf

- The term frequency $tf_{t,d}$ of term t in document d is defined as the number of times that t occurs in d .
- df_t is the document frequency of t : the number of documents that contain t
 - df_t is an inverse measure of the informativeness of t
 - $df_t \leq N$
- We define the idf (inverse document frequency) of t by
$$idf_t = \log_{10}(N/df_t)$$
 - We use $\log(N/df_t)$ instead of N/df_t to “dampen” the effect of idf.

Tf-idf

- The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$w_{t,d} = \log(1 + tf_{t,d}) \times \log_{10}(N / df_t)$$

- Best known weighting scheme in information retrieval
 - Note: the “-” in tf-idf is a hyphen, not a minus sign!
 - Alternative names: tf.idf, tf x idf
- Increases with the number of occurrences within a document
- Increases with the rarity of the term in the collection

Filtering for Cleaner Features

- **Stopwords**
 - weeding out common words that make for vacuous features
- **Frequency-Based Filtering**
 - filtering out corpus-specific common words as well as general-purpose stopwords
- **Rare words**
 - Depending on the task, one might also need to filter out rare words.
 - These might be truly obscure words, or misspellings of common words.
- **Stemming**
 - An NLP task that tries to chop each word down to its basic linguistic word stem form

Word representation: embedding the context

- Attempt to encode similarity inside the word vectors
- Built ontop of the following great idea
 - “You shall know a word by the company it keeps” (J. R. Firth 1957)

During his presidency, **Trump** ordered a travel ban on citizens controversial or false. **Trump** was elected president in a surprise victory over 1971, renamed it to The **Trump** Organization, and expanded it into Manhattan. coordination between the **Trump** campaign and the Russian government in its election interference.



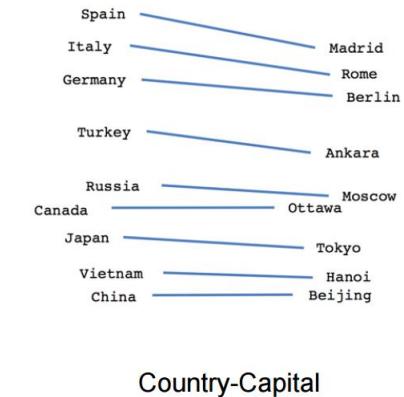
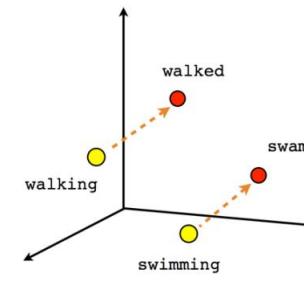
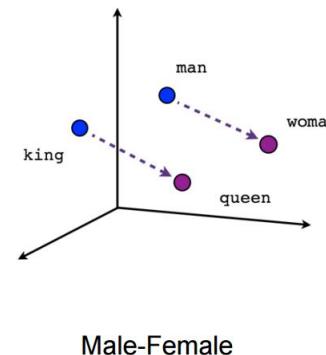
These words describe the meaning of Trump

Word embedding

- Each word is encoded in a dense vector (Low dimension)
- Able to capture the semantics
 - Similar words ~ Similar vectors

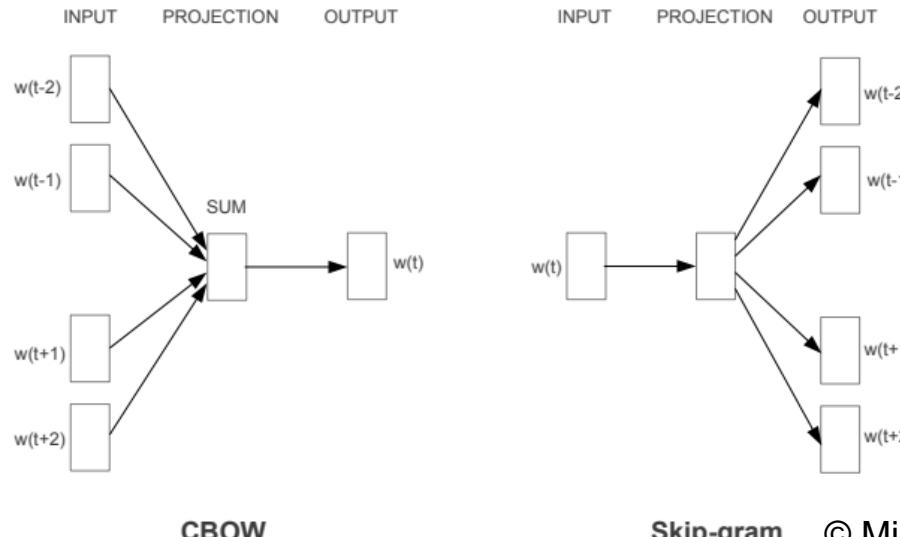
University =

[0.13
0.67
-
0.34
0.76
-
0.21
-0.11
-
0.45
0.87
0.44]



How to learn word embeddings

- The famous approach: Word2vec (Mikolov et. al. 2013)
- Unsupervised learning
- Large-scale dataset
- Lower computation cost
- **High quality word vectors**



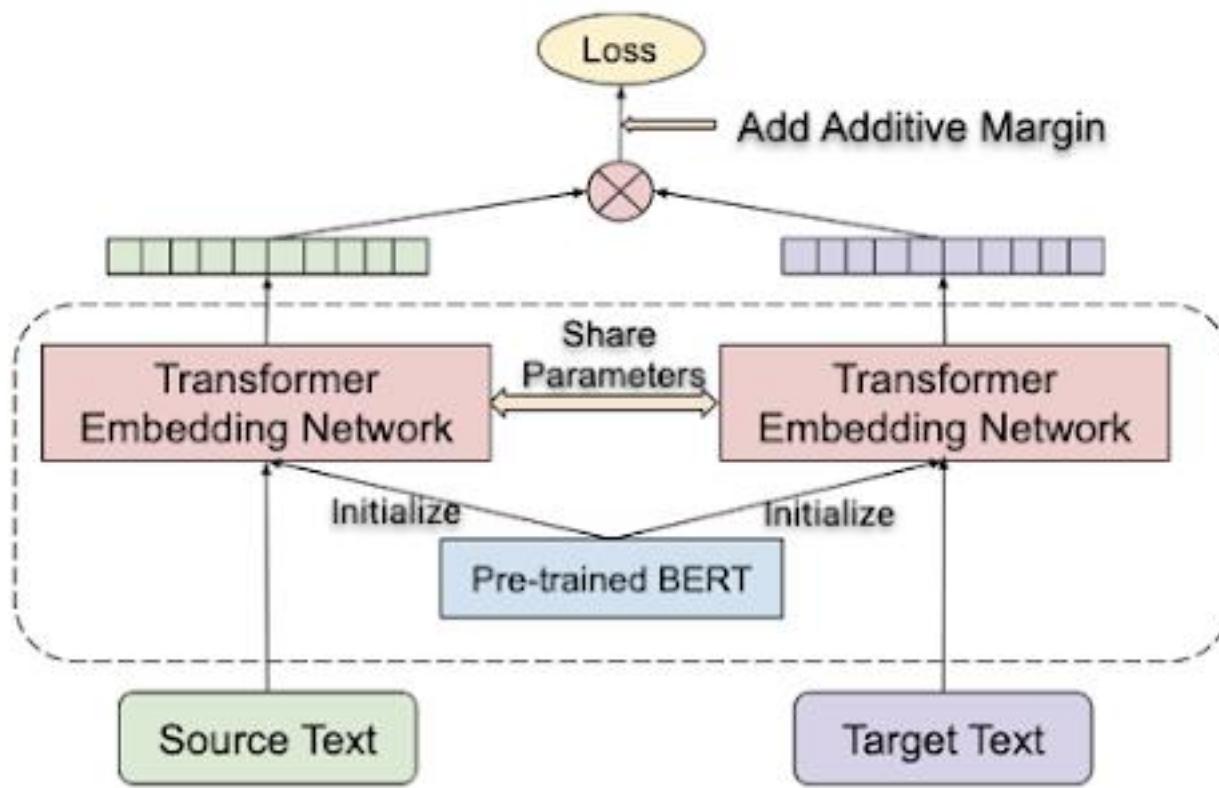
CBOW

Skip-gram

© Mikolov et. Al. 2013

BERT sentence embedding

- Feng, Fangxiaoyu, et al. "Language-agnostic BERT Sentence Embedding." *arXiv preprint arXiv:2007.01852* (2020).



Feature selection

Interaction Features

- A simple pairwise *interaction feature* is the product of two features
- A simple linear model
 - $y = w_1x_1 + w_2x_2 + \dots + w_nx_n$
- An easy way to extend the linear model is to include combinations of pairs of input features
 - $y = w_1x_1 + w_2x_2 + \dots + w_nx_n + w_{1,1}x_1x_1 + w_{1,2}x_1x_2 + w_{1,3}x_1x_3 + \dots$

Polynomial Features

$$(X_1, X_2) \longrightarrow (1, X_1, X_2, X_1^2, X_1X_2, X_2^2)$$

```
>>> import numpy as np
>>> from sklearn.preprocessing import PolynomialFeatures
>>> X = np.arange(6).reshape(3, 2)
>>> X
array([[0, 1],
       [2, 3],
       [4, 5]])
>>> poly = poly = PolynomialFeatures(degree=2, interaction_only=False,
include_bias=True)
>>> poly.fit_transform(X)
array([[ 1.,  0.,  1.,  0.,  0.,  1.],
       [ 1.,  2.,  3.,  4.,  6.,  9.],
       [ 1.,  4.,  5., 16., 20., 25.]])
```

Polynomial features with scikit-learn

Feature Selection

- Objective
 - prune away nonuseful features in order to reduce the complexity of the resulting model
- Advantages
 - **Training** a machine learning algorithm **faster**.
 - Reducing the **complexity** of a model and making it easier to **interpret**.
 - Building a **sensible model** with **better prediction power**.
 - **Reducing overfitting** by selecting the right set of features.

Wrapper methods

- The *feature selection* process is based on a specific machine learning algorithm
- Exhaustive search follows a *greedy search approach* by evaluating all the possible combinations of features against the *evaluation criterion*
- Random search methods randomly generate a subset of features
- Computationally intensive since for each subset a new model needs to be trained

Initial set of
all features

Embedded Methods

- Perform feature selection during the model training
- Decision tree
 - select a feature in each recursive step of the tree growth process and divide the sample set into smaller subsets
 - The more child nodes in a subset are in the same class, the more informative the features are

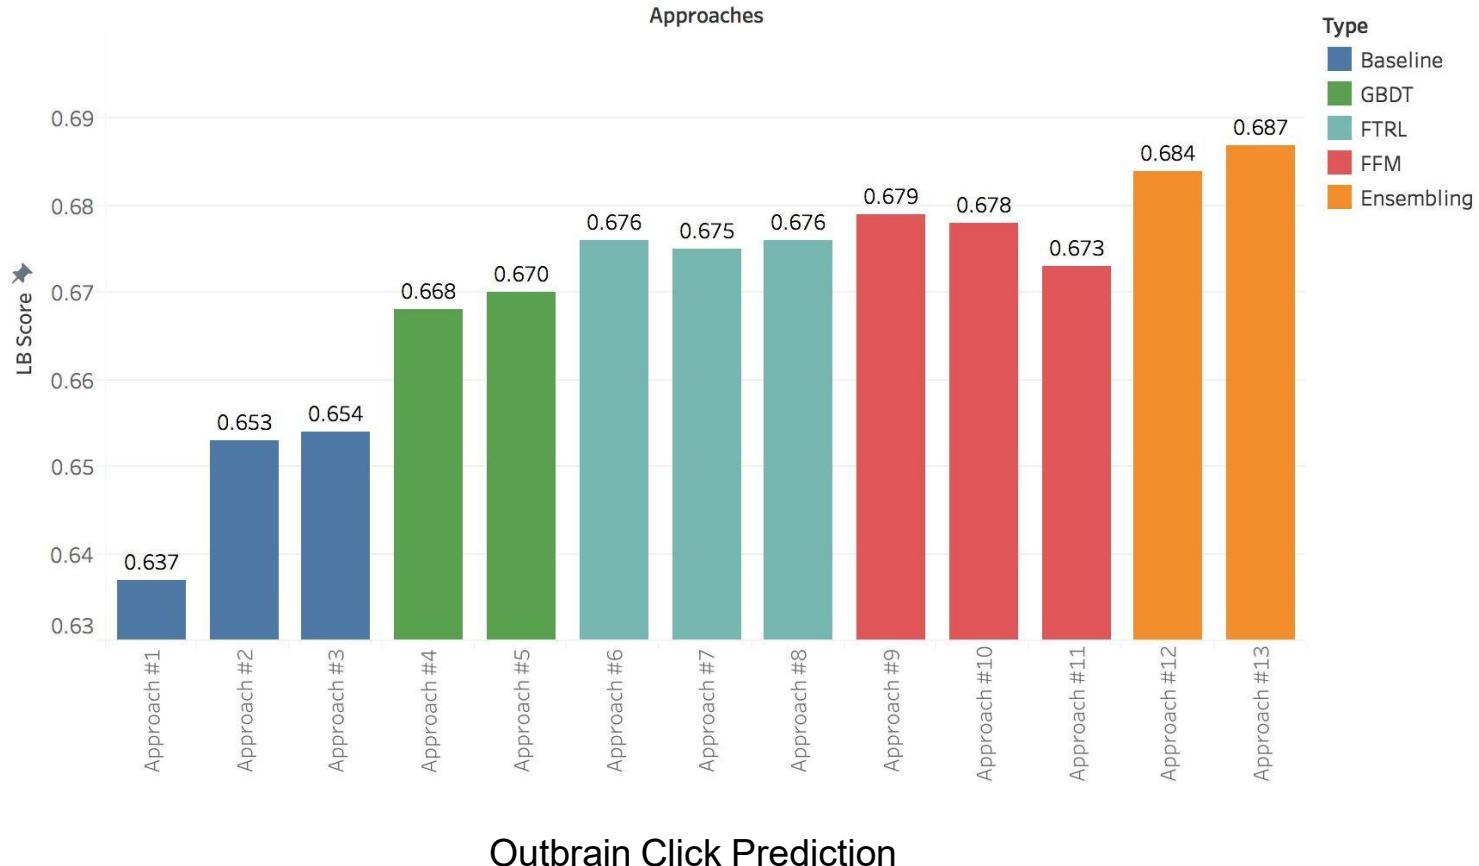
Method comparision

Filter methods	Wrapper methods	Embedded methods
Generic set of methods which do not incorporate a specific machine learning algorithm.	Evaluates on a specific machine learning algorithm to find optimal features.	Embeds (fix) features during model building process. Feature selection is done by observing each iteration of model training phase.
Much faster compared to Wrapper methods in terms of time complexity	High computation time for a dataset with many features	Sits between Filter methods and Wrapper methods in terms of time complexity
Less prone to over-fitting	High chances of over-fitting because it involves training of machine learning models with different combination of features	Generally used to reduce over-fitting by penalizing the coefficients of a model being too large.
Examples – Correlation, Chi-Square test, ANOVA, Information gain etc.	Examples - Forward Selection, Backward elimination, Stepwise selection etc.	Examples - LASSO, Elastic Net, Ridge Regression etc.

“More data beats clever algorithms, but better data beats more data.”

— Peter Norvig

Diverse set of features and models leads to different results!



Towards Automated Feature Engineering

Deep Learning....



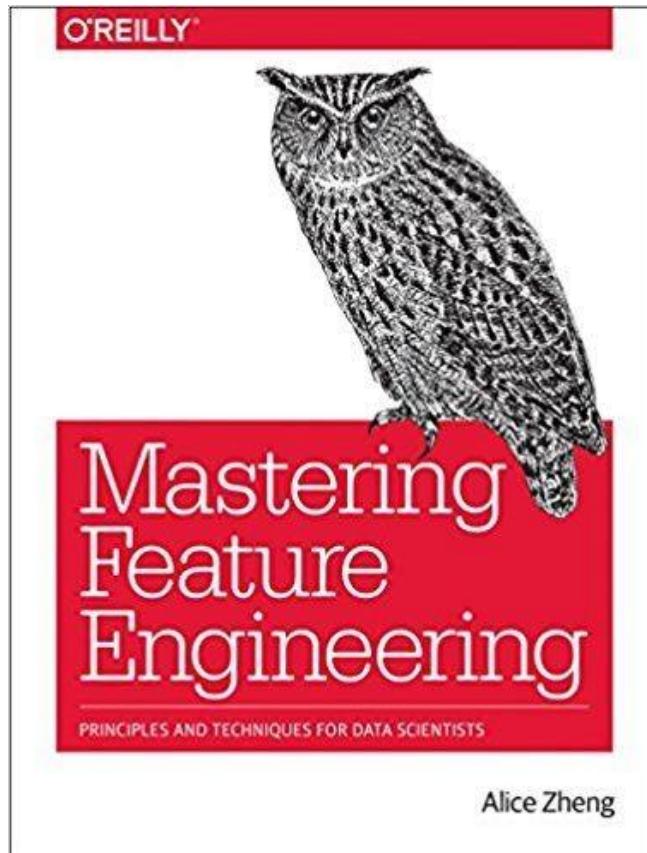
25
YEARS ANNIVERSARY
SOICT

VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Thank you
for your
attention!!!



References



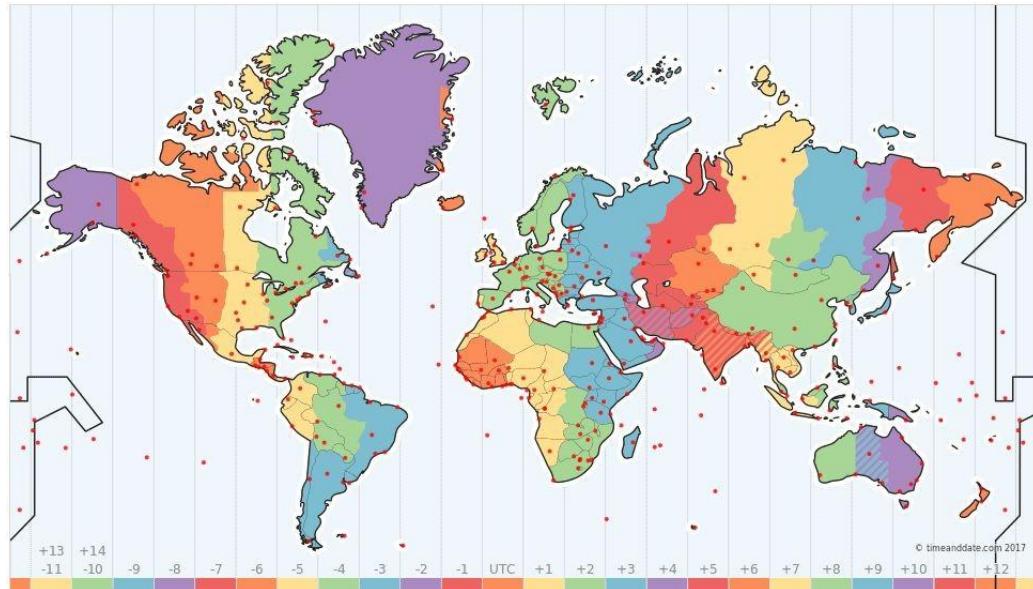
- Scikit-learn - Preprocessing data
- Spark ML - Feature extraction

Temporal Features



Time Zone conversion

- Factors to consider:
 - Multiple time zones in some countries
 - Daylight Saving Time (DST)
 - Start and end DST dates



	country_name	utc_time_offset	dst_time_offset
0	Afghanistan	+04:30	-
1	Aaland Islands	+02:00	+03:00
2	Albania	+01:00	+02:00
3	Algeria	+01:00	-
4	Samoa (American)	-11:00	-
5	Andorra	+01:00	+02:00
6	Angola	+01:00	-
7	Anguilla (UK)	-04:00	-
8	Antigua & Barbuda	-04:00	-
9	Argentina	-03:00	-

Time binning

- Apply binning on time data to make it categorial and more general.
- Binning a time in hours or periods of day, like below.

Hour range	Bin ID	Bin Description
[5, 8)	1	Early Morning
[8, 11)	2	Morning
[11, 14)	3	Midday
[14, 19)	4	Afternoon
[19, 22)	5	Evening
[22-24) and (00-05]	6	Night

- Extraction: weekday/weekend, weeks, months, quarters, years...

Closeness to major events

- Hardcode categorical features from dates
- Example: Factors that might have major influence on spending behavior
- Proximity to major events (holidays, major sports events)
 - Eg. date_X_days_before_holidays
- Proximity to wages payment date (monthly seasonality)
 - Eg. first_saturday_of_the_month



Time differences

- Differences between dates might be relevant
- Examples:
 - user_interaction_date - published_doc_date
- To model how recent was the ad when the user viewed it. Hypothesis: user interests on a topic may decay over time
 - last_user_interaction_date - user_interaction_date
- To model how old was a given user interaction compared to his last interaction

Spatial Features

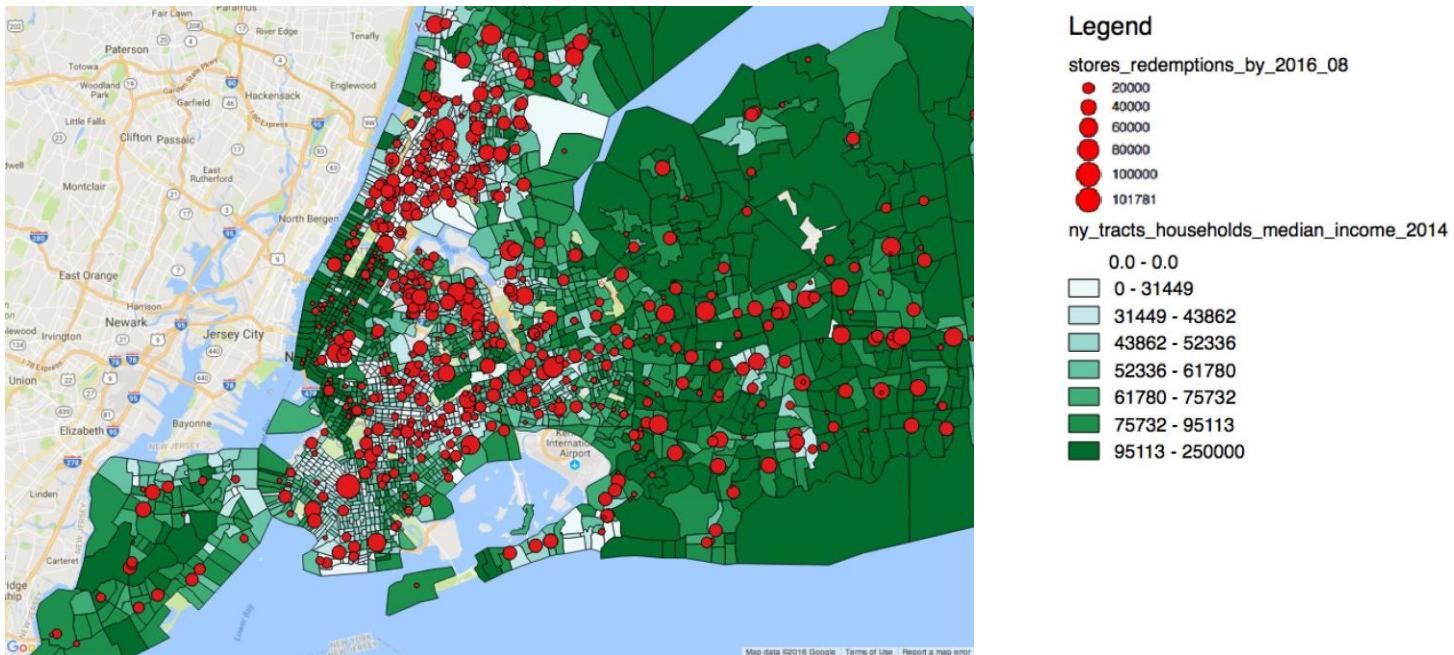


Spatial Variables

- Spatial variables encode a location in space, like:
 - GPS-coordinates (lat. / long.) - sometimes require projection to a different [coordinate system](#)
 - Street Addresses - require geocoding
 - ZipCodes, Cities, States, Countries - usually enriched with the centroid coordinate of the polygon (from external GIS data)
- Derived features
 - Distance between a user location and searched hotels ([Expedia competition](#))
 - Impossible travel speed (fraud detection)

Spatial Enrichment

- Usually useful to enrich with external geographic data (eg. Census demographics)



Beverage Containers Redemption Fraud Detection: Usage of # containers redeemed (red circles) by store and Census households median income by Census Tracts